


Programowanie w Pythonie

Łukasz Mioduszewski, UKSW 2022

Wstęp do sieci neuronowych



Sieci Neuronowe - historia

- 1933 - Edward Thorndike: human learning consists in the strengthening of some (then unknown) property of neurons
- 1949 - Donald Hebb: it is specifically a strengthening of *connections* between neurons in the brain that accounts for learning

„When an axon of cell A is near enough to excite cell B and repeatedly or persistently takes part in firing it, some growth process or metabolic change takes place in one or both cells such that A's efficiency, as one of the cells firing B, is increased.”


D. Hebb, The Organization of Behavior, 1949

"Neurons that fire together wire together."

(Hebb's Law)


Sieć Hopfielda

- Pomysł Hopfielda (1982): jeden „neuron” odpowiada jednemu pikselowi obrazu (czarno-białego, a więc ma wartość +1 albo -1)
- Każdy neuron może być „połączony” z każdym przy użyciu macierzy wagi \mathbf{w}
- Dążymy do minimalizacji energii




$L=5$

$M=L \times L$

$x_i = \pm 1$ 

$$E = -\frac{1}{2} \sum_{ij} w_{ij} x_i x_j$$



Uzupełnienie matematyczne

- Iloczyn zewnętrzny

$$\begin{bmatrix} 1 \\ -1 \\ -1 \end{bmatrix} \begin{bmatrix} 1 & -1 & -1 \end{bmatrix} = \begin{bmatrix} 0 & -1 & -1 \\ -1 & 0 & 1 \\ -1 & 1 & 0 \end{bmatrix}$$

- Macierz jednostkowa

$$I = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$



Sieć Hopfielda

- „uczymy” sieć nowych wzorów, może ich być n :



"Neurons that fire together, wire together. Neurons that fire out of sync, fail to link".

$$\mathbf{w} = \frac{1}{n} \sum_{\mu=1}^n (\mathbf{x}^{\mu} \otimes \mathbf{x}^{\mu} - \mathbf{1})$$

To są macierze!

- Takie same wartości dadzą dodatnie w , czyli ujemną energię

$x_i = \pm 1$


$$E = -\frac{1}{2} \sum_{ij} w_{ij} x_i x_j$$


Sieć Hopfielda

- energia i-tego spinu:

$$E_i = -\frac{1}{2}x_i \left(\sum_j w_{ij} x_j \right) = -\frac{1}{2}x_i h(i)$$

- Ewolucja w czasie:

$$\mathbf{x}(t+1) = \text{sgn}(\mathbf{w} \cdot \mathbf{x}(t))$$

$x_i = \pm 1$



$$E = -\frac{1}{2} \sum_{ij} w_{ij} x_i x_j$$


Sieć Hopfielda w praktyce

- Zadanie: prosty OCR
(Optical Character Recognition)

A = ""	Z = ""
OXXXO	XXXXX
XOOOX	OOOXO
XXXXX	OOXOO
XOOOX	OXOOO
XOOOX	XXXXX
""	""

- Wprowadzamy dwa wzory do nauczania się: litery A i Z

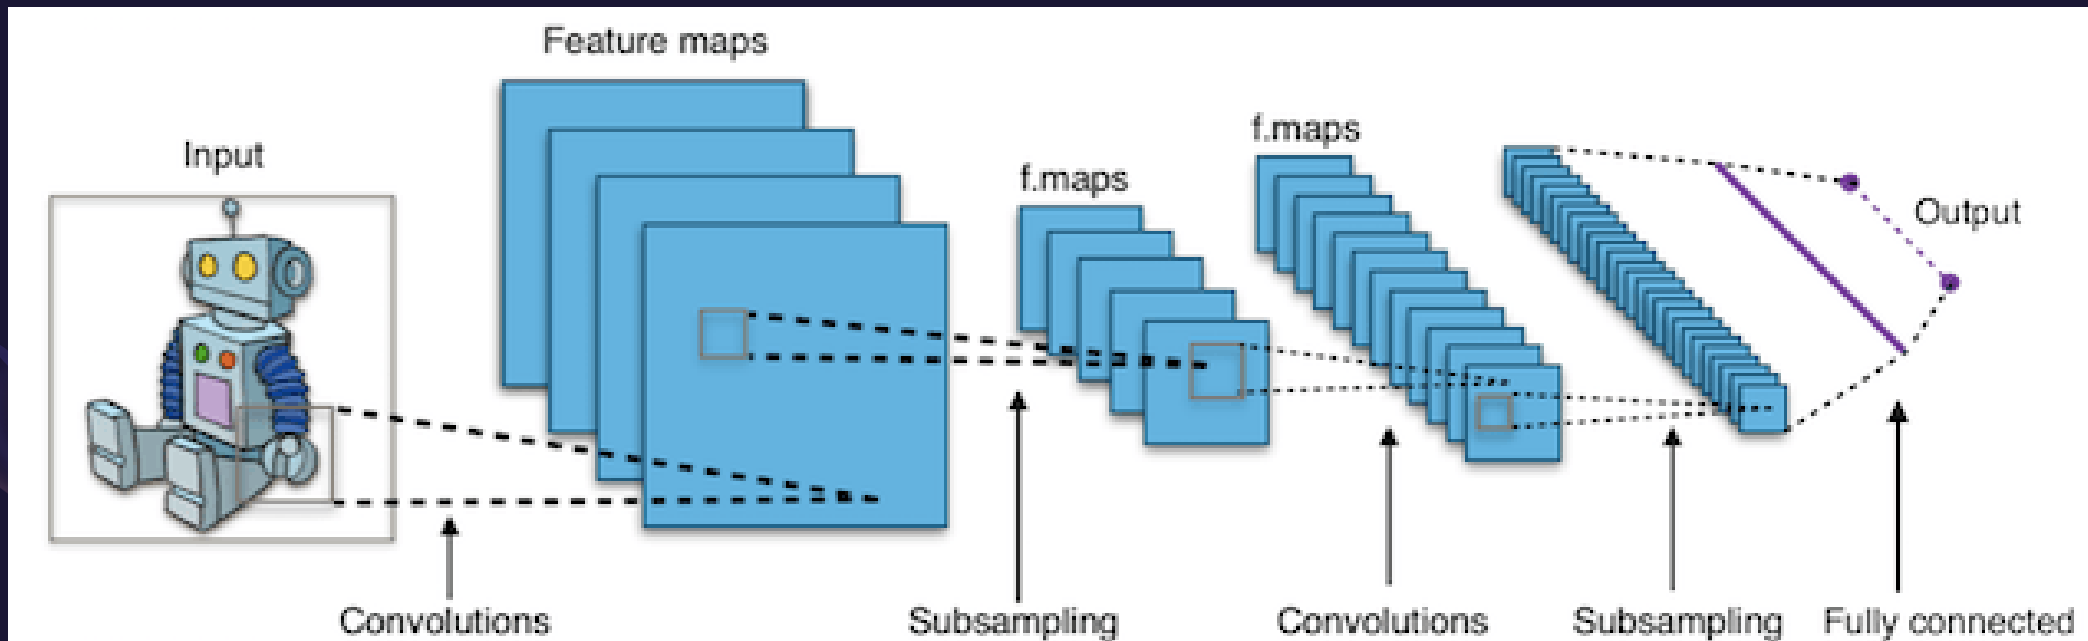
- Reprezentujemy je jako obrazki 5x5 pikseli



- Zaczynamy z czegoś zdeformowanego

Sieci neuronowe dziś

- Każdy neuron pobiera **dane** (liczby), każdą daną mnoży przez **wagę**, czego wynikiem jest liczba, która staje się daną dla kolejnej warstwy neuronów
- **Deep** learning – na tyle wiele warstw że nie wiemy co się dzieje



Biblioteka Keras

- Wysokopoziomowy "frontend" do TensorFlow od Google
- Alternatywa: FastAI ("frontend" do PyTorch od Facebook/Meta)
- Oba są open source



Biblioteka Keras

- Wymagania:
 - Python 3+
 - SciPy z NumPy
 - Matplotlib (tylko do rysowania, niekonieczny)
- Instalacja przez pip:

```
pip install tensorflow # zwykle wystarczy
pip install keras      # jeśli nie, to jeszcze to
```
- Instalacja przez Anacondę: <https://elitedatascience.com/keras-tutorial-deep-learning-in-python>

Początek skryptu

keras_CNN_example.py

```
1 import numpy as np
2 np.random.seed(123) # for reproducibility
3 from keras.models import Sequential # sequential neural network
4 from keras.layers import Dense, Dropout, Activation, Flatten # usual layers
5 from keras.layers import Convolution2D, MaxPooling2D # to train on image data
6 from keras.utils import np_utils # some utilities
7
8
9
10
```

Ładowanie bazy danych do uczenia sieci

keras_CNN_example.py

```
1 import numpy as np
2 np.random.seed(123) # for reproducibility
3 from keras.models import Sequential # sequential neural network
4 from keras.layers import Dense, Dropout, Activation, Flatten # usual layers
5 from keras.layers import Convolution2D, MaxPooling2D # to train on image data
6 from keras.utils import np_utils # some utilities
7
8 from keras.datasets import mnist
9 # Load pre-shuffled MNIST data into train and test sets
10 (X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Ładowanie bazy danych do uczenia sieci

keras_CNN_example.py


```
1 import numpy as np
2 np.random.seed(123) # for reproducibility
3 from keras.models import Sequential # sequential neural network
4 from keras.layers import Dense, Dropout, Activation, Flatten # usual layers
5 from keras.layers import Convolution2D, MaxPooling2D # to train on image data
6 from keras.utils import np_utils # some utilities
7
8 from keras.datasets import mnist
9 # Load pre-shuffled MNIST data into train and test sets
10 (X_train, y_train), (X_test, y_test) = mnist.load_data()
11 print( X_train.shape ) # 60k samples, each is a 28x28 black-and-white image
11 # (60000, 28, 28)

# X_train to dane do treningu, y_train to poprawne wyniki dla tych danych
# X_test to dane testujące sieć, y_test to pożądaný wynik
```

Ładowanie bazy danych do uczenia sieci

keras_CNN_example.py

```
1 import numpy as np
2 np.random.seed(123) # for reproducibility
3 from keras.models import Sequential # sequential neural network
4 from keras.layers import Dense, Dropout, Activation, Flatten # usual layers
5 from keras.layers import Convolution2D, MaxPooling2D # to train on image data
6 from keras.utils import np_utils # some utilities
7
8 from keras.datasets import mnist
9 # Load pre-shuffled MNIST data into train and test sets
10 (X_train, y_train), (X_test, y_test) = mnist.load_data()
11 print( X_train.shape ) # 60k samples, each is a 28x28 black-and-white image
11 # (60000, 28, 28)
12 from matplotlib import pyplot as plt
13 plt.imshow(X_train[0])
```



Preprocessing inputu

keras_CNN_example.py

```
14 # tensorflow must know how many channes there are, so the last dimension
15 should be the number of channels (in this case, one)
16 X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
17 X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)
18 print( X_train.shape )
19 # (60000, 28, 28, 1)
20 X_train = X_train.astype('float32') # input data must be float32
21 X_test = X_test.astype('float32')
22 X_train /= 255 # input data must be in range from 0 to 1
23 X_test /= 255
24
25
26
27
28
29
30
```


Preprocessing outputu

keras_CNN_example.py

```
14 # tensorflow must know how many channes there are, so the last dimension
15 should be the number of channels (in this case, one)
16 X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
17 X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)
18 print( X_train.shape )
19 # (60000, 28, 28, 1)
20 X_train = X_train.astype('float32') # input data must be float32
21 X_test = X_test.astype('float32')
22 X_train /= 255 # input data must be in range from 0 to 1
23 X_test /= 255
24
25 print( y_train[:10] )
26 # [5 0 4 1 9 2 1 3 1 4] # y_train tells which digit is showed in every picture
27 # Convert 1-dimensional class arrays to 10-dimensional class matrices
28 Y_train = np_utils.to_categorical(y_train, 10)
29 Y_test = np_utils.to_categorical(y_test, 10)
30 print( Y_train.shape ) # (60000, 10)
```

Konwolucja

- Konwolucja pomiędzy obrazkiem a filtrem/jądrem konwolucji (convolution kernel/filter, tak jak te okna Hanninga albo Blackmana)

$$g(x, y) = \omega * f(x, y) = \sum_{dx=-a}^a \sum_{dy=-b}^b \omega(dx, dy) f(x + dx, y + dy)$$

- g to wyjściowy obraz, f to wejściowy obraz, w to filtr

	0	1	2	3
0	25	80	110	5
1	80	100	120	7
2	92	60	30	25
3	93	97	23	6

Kernel (sharpen)

$$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$$



	0	1	2	3
0				
1		160		
2				
3				

Konwolucja

$$\begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$



$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$



Name	Kernel	Image Result
Identity	$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	
Sharpen	$\begin{bmatrix} 0 & -1 & 0 \\ -1 & 5 & -1 \\ 0 & -1 & 0 \end{bmatrix}$	
Mean Blur	$\begin{bmatrix} 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \\ 1/9 & 1/9 & 1/9 \end{bmatrix}$	
Laplacian	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	
Gaussian Blur	$\begin{bmatrix} 1/16 & 2/16 & 1/16 \\ 2/16 & 4/16 & 2/16 \\ 1/16 & 2/16 & 1/16 \end{bmatrix}$	

Architektura sieci

keras_CNN_example.py

```
32 model = Sequential()
33
34 # defining input layer. input_shape should be the shape of 1 sample.
35 # here it's (28, 28, 1) = (width, height, channels) of each digit image.
36 model.add(Convolution2D(32, (3,3), activation='relu', input_shape=(28,28,1)))
37 # 32 convolution filters, each of them is a 3x3 matrix
38 print( model.output_shape )
39 # (None, 26, 26, 32) corresponds to (samples, new_rows, new_cols, filters)
40 # model will output all samples, convoluted into 26x26 array using 32 filters
41
42 # adding more layers
43 model.add(Convolution2D(32, (3,3), activation='relu')) # regularization max(x,0)
44
45 model.add(MaxPooling2D(pool_size=(2,2))) # way to reduce number of parameters
46 in our model by sliding a 2x2 pooling filter across the previous layer and
47 taking the max of the 4 values in the 2x2 filter
48 model.add(Dropout(0.25)) # killing some neurons
```

Architektura sieci – idziemy głębiej

keras_CNN_example.py

```
50 model.add(Flatten()) # flattens the input (1 channel)
51 model.add(Dense(128, activation='relu')) # 128 is the output size
52 # keras automatically matches layer input/output sizes
53 model.add(Dropout(0.5)) # killing excessive neurons again
54 model.add(Dense(10, activation='softmax')) # normalize the output
55 # to a probability distribution over predicted output classes
56
57
58
59
60
61
62
63
64
65
```

Kompilacja modelu

keras_CNN_example.py

```
50 model.add(Flatten()) # flattens the input (1 channel)
51 model.add(Dense(128, activation='relu')) # 128 is the output size
52 # keras automatically matches layer input/output sizes
53 model.add(Dropout(0.5)) # killing excessive neurons again
54 model.add(Dense(10, activation='softmax')) # normalize the output
55 # to a probability distribution over predicted output classes
56
57 model.compile(loss='categorical_crossentropy', optimizer='adam',
58               metrics=['accuracy'])
59 # loss function defines the "cost" associated with good/bad image recognition
60
61
62
63
64
65
```

Trenujemy nasz model

keras_CNN_example.py

```
50 model.add(Flatten()) # flattens the input (1 channel)
51 model.add(Dense(128, activation='relu')) # 128 is the output size
52 # keras automatically matches layer input/output sizes
53 model.add(Dropout(0.5)) # killing excessive neurons again
54 model.add(Dense(10, activation='softmax')) # normalize the output
55 # to a probability distribution over predicted output classes
56
57 model.compile(loss='categorical_crossentropy', optimizer='adam',
58               metrics=['accuracy'])
59 # loss function defines the "cost" associated with good/bad image recognition
60
61 model.fit(X_train, Y_train,
62           batch_size=32, epochs=10, verbose=1)
63 # Epoch 1/10
64 # 7744/60000 [==>.....] - ETA: 96s - loss: 0.5806 - acc: 0.816
65
```


Testujemy nasz model

keras_CNN_example.py

```
50 model.add(Flatten()) # flattens the input (1 channel)
51 model.add(Dense(128, activation='relu')) # 128 is the output size
52 # keras automatically matches layer input/output sizes
53 model.add(Dropout(0.5)) # killing excessive neurons again
54 model.add(Dense(10, activation='softmax')) # normalize the output
55 # to a probability distribution over predicted output classes
56
57 model.compile(loss='categorical_crossentropy', optimizer='adam',
58               metrics=['accuracy'])
59 # loss function defines the "cost" associated with good/bad image recognition
60
61 model.fit(X_train, Y_train,
62           batch_size=32, epochs=10, verbose=1)
63 # Epoch 10/10
64
65 score = model.evaluate(X_test, Y_test, verbose=0)
```

A jak potem tego używać?

- `y_krm = model.predict(x)`
- Można też używać metody `call()`, ale są pewne subtelne różnice:
https://stackoverflow.com/questions/60837962/confusion-about-keras-model-call-vs-call-vs-predict-methods?fbclid=IwAR0wAIRHfSa2o6-qriOgLp64lfXhNtLr8kk8f9jF9Q67_hdgJVvjzM5MBW8
- Aby zapisać model:

```
model = ... # Get model (Sequential, Functional Model, or Model subclass)
model.save('path/to/location')
```
- Aby odczytać model:

```
from tensorflow import keras
model = keras.models.load_model('path/to/location')
```

Bibliografia

- <https://elitedatascience.com/keras-tutorial-deep-learning-in-python>
- <https://www.knowledgehut.com/blog/data-science/pytorch-vs-tensorflow>
- <https://medium.com/@ianormy/convolution-filters-4971820e851f>
- <https://www.activestate.com/resources/quick-reads/what-is-a-keras-model/>

